

MicroHack challenge 2

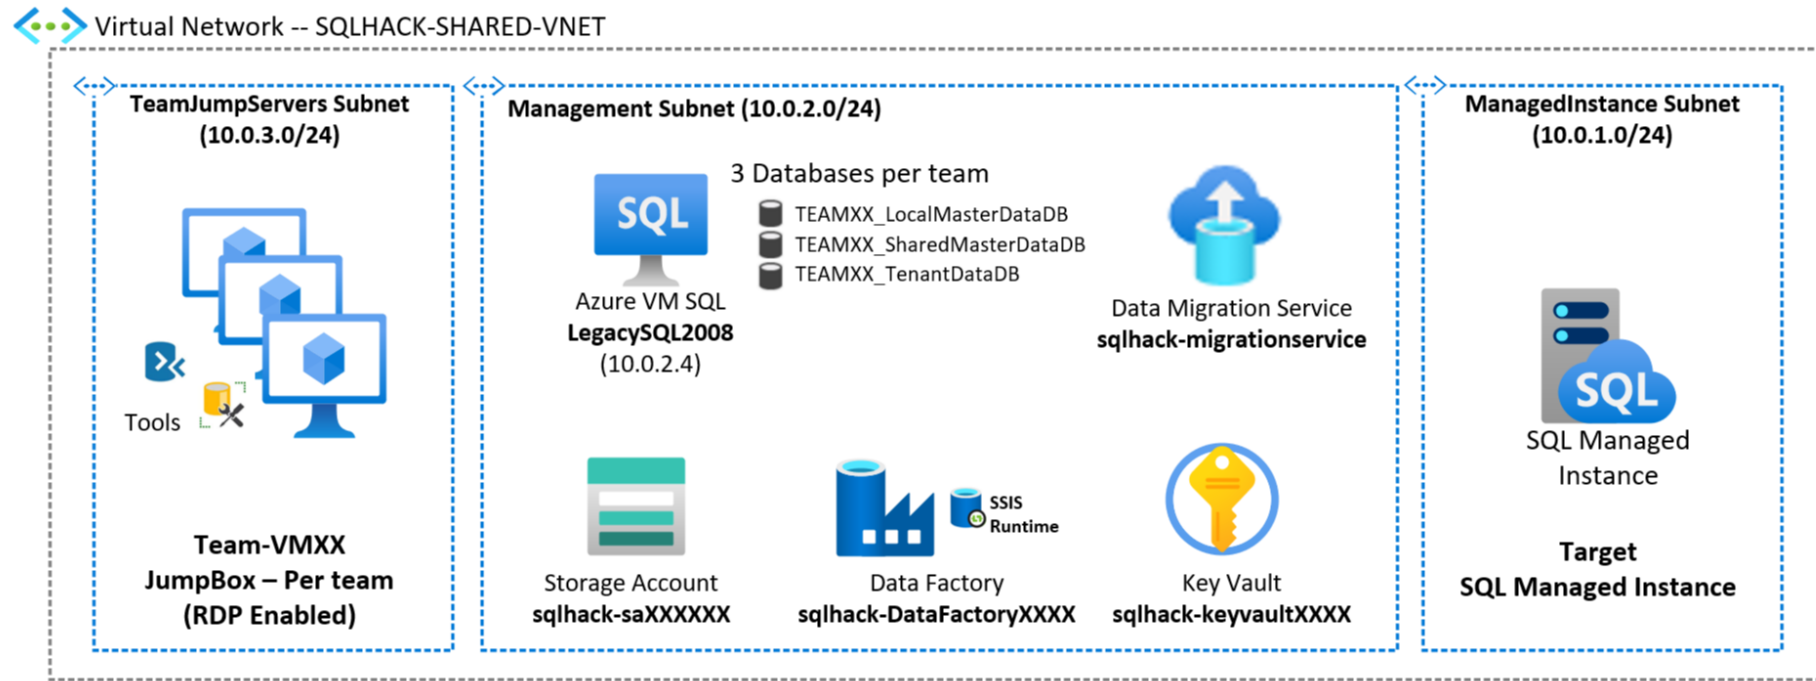
Monitoring and Performance on Azure SQL Managed Instance

V4.0

Contents

Migration architecture and Azure components	2
Lab Overview & Background	3
Background	3
LAB 1: Identifying performance issues, their root causes and fixing the problem	4
1. Using DMVs and the Query Store to identify performance bottlenecks	4
2. Using the Azure Portal to identify performance bottlenecks	10
3. Using Log Analytics with KQL to visualize SQL MI and Database Stats	12

Migration architecture and Azure components



SQLHACK-SHARED-VNET		
Single Virtual Network containing all workshop resources		
TeamJumpServers Subnet Each team is assigned a Win10 VM that mimics their company desktop	Management Subnet Several machines and services are already deployed within a dedicated subnet within the Virtual Network	ManagedInstance Subnet The Azure SQL Managed Instance has been deployed into a dedicated Subnet

Lab Overview & Background

In this exercise you will explore how you can monitor and diagnose performance issues with Azure SQL Database Managed Instance using SQL Server Management Studio, DMVs (Dynamic Management Views) and Azure Portal tools.

Background

We have seen that the legacy application is a multi-tenant system with a collection of 3 databases supporting the transactional workload for each customer – these are the 3 databases we have modernised by migrating to the shared Azure SQL Managed Instance.

Apart from the tenant centric transactional databases there are 2 shared databases:

- **2008DW:** A centralised Data Warehouse database that combines data from the various tenant transactional databases and is used for aggregated reporting and analytics.
- **TenantCRM:** A centralised database that is used to manage all customer relationships and order processing.

Recently you have heard complaints from users that the TenantCRM system is running slowly. Recently changes were made to some of the stored procedures in the TenantCRM database and the App dev team believe these are the source of the problem. They've asked for your help to track down the performance issues and its root cause.

Let's follow the steps below to begin to troubleshoot this issue.

The exercises in this lab are conducted against the shared [TenantCRM] database

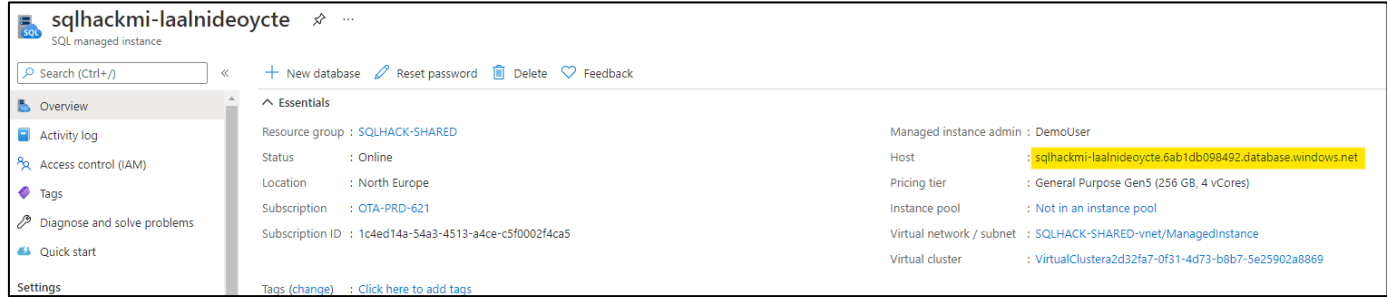
LAB 1: Identifying performance issues, their root causes and fixing the problem

1. Using DMVs and the Query Store to identify performance bottlenecks.

Performance issues in SQL Server can be grouped into 1 of 2 high-level categories: *Running* or *Waiting*.

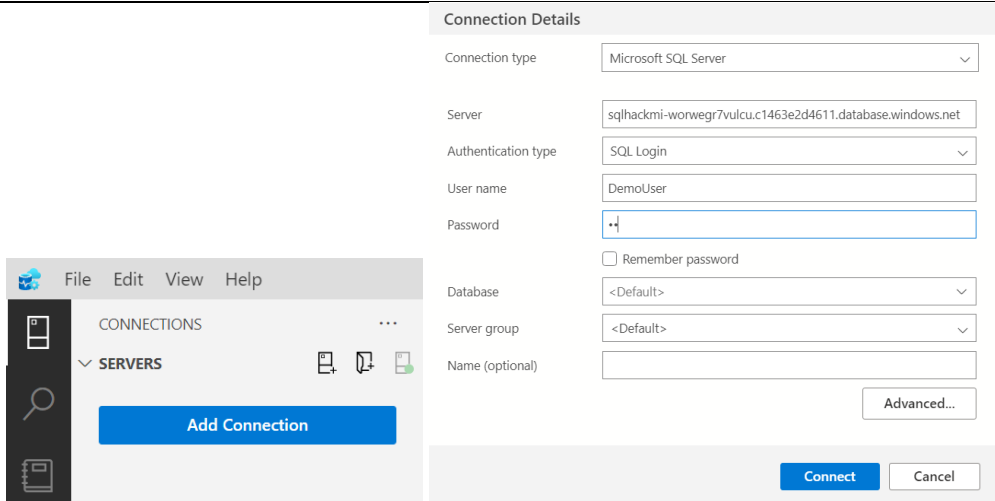
Running means a query has been allocated CPU resource and is actively processing but is taking a long time to complete. Waiting means a query is not actively being processed by the CPU but is waiting on another resource (such as memory, disk, network, etc) to pass data in so processing can continue.

Finding the bottleneck will help us determine what category issue we have so we can progress to identifying the root cause. Let's see how to find the bottleneck using both SSMS and the Azure portal.

Instructions/Narrative	Screenshot	Notes
<p>On your Win10 Team VM open the Azure Portal and navigate to the shared SQL Managed Instance.</p> <p>Scroll down the Overview screen copy the host name of the Managed Instance and save it somewhere handy.</p>		

On your Team Win10 VM open **Azure Data Studio** and **Add Connection** using these details:

Server: <just copied>
Authentication: **SQL Login**
Login:
Will be provided during hack
Password:
Will be provided during hack

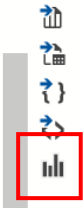
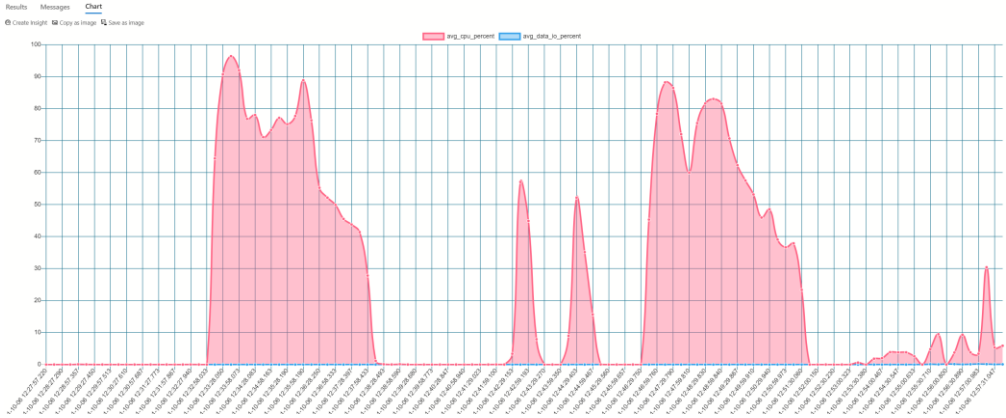


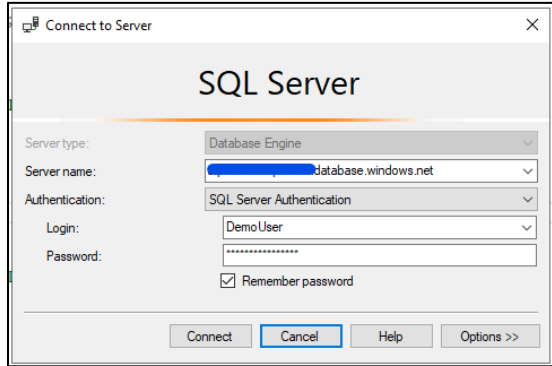
The screenshot shows the Azure Data Studio interface. On the left, the 'CONNECTIONS' pane is visible with a blue 'Add Connection' button. On the right, the 'Connection Details' dialog box is open, showing the following fields:

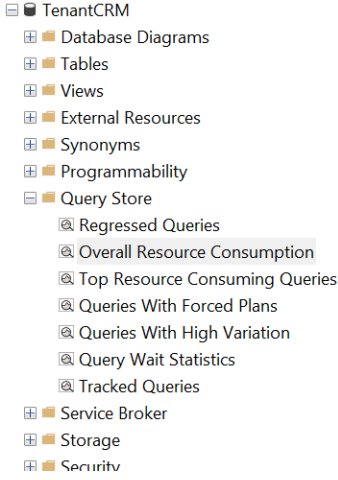
- Connection type: Microsoft SQL Server
- Server: sqlhackmi-worwegr7vulcu.c1463e2d4611.database.windows.net
- Authentication type: SQL Login
- User name: DemoUser
- Password: (masked with dots)
- Remember password: ☐
- Database: <Default>
- Server group: <Default>
- Name (optional):


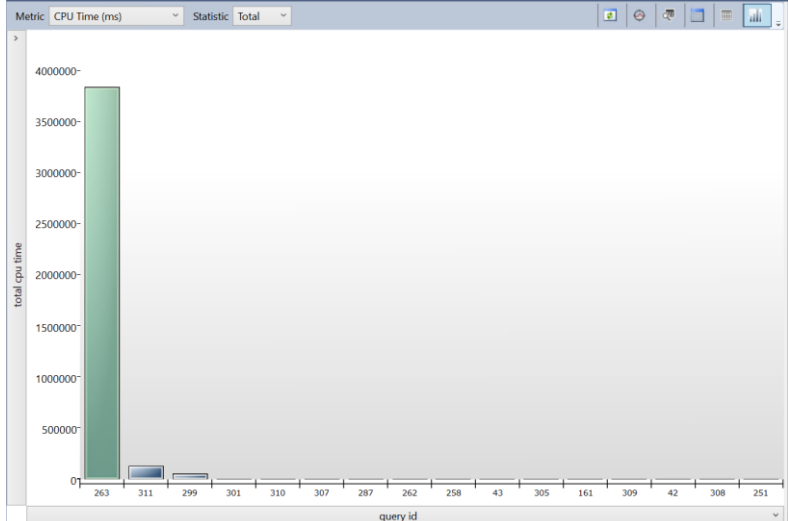
At the bottom of the dialog, there are 'Connect' and 'Cancel' buttons.

Narrative/Instructions	Screenshot	Notes																																																												
<p>Open the SQL script: C:_SQLHACK_\LABS\02-Administering_Monitoring\Part_01_Monitoring_Lab_1.sql</p> <p>Run the PART 1 query 1.</p> <p>Look at the results and note the [avg_cpu_percent] value is about 70% indicating that this database is consuming a lot of the hosts CPU resources</p>	<table><tr><th colspan="5"></th></tr><tr><th colspan="2">Results</th><th colspan="3">Messages</th></tr><tr><th></th><th>end_time</th><th>avg_cpu_percent</th><th>avg_data_io_percent</th><th>avg_log_write_percent</th></tr><tr><td>1</td><td>2021-10-06 13:03:31.780</td><td>81.17</td><td>0.00</td><td>0.00</td></tr><tr><td>2</td><td>2021-10-06 13:03:16.770</td><td>77.06</td><td>0.00</td><td>0.00</td></tr><tr><td>3</td><td>2021-10-06 13:03:01.770</td><td>42.29</td><td>0.00</td><td>0.00</td></tr><tr><td>4</td><td>2021-10-06 13:02:46.747</td><td>0.00</td><td>0.00</td><td>0.00</td></tr><tr><td>5</td><td>2021-10-06 13:02:31.727</td><td>0.00</td><td>0.00</td><td>0.00</td></tr><tr><td>6</td><td>2021-10-06 13:02:16.680</td><td>0.00</td><td>0.00</td><td>0.00</td></tr><tr><td>7</td><td>2021-10-06 13:02:01.633</td><td>0.00</td><td>0.96</td><td>0.00</td></tr><tr><td>8</td><td>2021-10-06 13:01:46.600</td><td>0.00</td><td>0.00</td><td>0.00</td></tr><tr><td>9</td><td>2021-10-06 13:01:31.570</td><td>0.00</td><td>0.00</td><td>0.00</td></tr></table>						Results		Messages				end_time	avg_cpu_percent	avg_data_io_percent	avg_log_write_percent	1	2021-10-06 13:03:31.780	81.17	0.00	0.00	2	2021-10-06 13:03:16.770	77.06	0.00	0.00	3	2021-10-06 13:03:01.770	42.29	0.00	0.00	4	2021-10-06 13:02:46.747	0.00	0.00	0.00	5	2021-10-06 13:02:31.727	0.00	0.00	0.00	6	2021-10-06 13:02:16.680	0.00	0.00	0.00	7	2021-10-06 13:02:01.633	0.00	0.96	0.00	8	2021-10-06 13:01:46.600	0.00	0.00	0.00	9	2021-10-06 13:01:31.570	0.00	0.00	0.00	<p>[sys].[dm_db_resource_stats] exposes a number of key database performance metrics.</p> <p>One metric every 15 seconds over the last hour</p>
Results		Messages																																																												
	end_time	avg_cpu_percent	avg_data_io_percent	avg_log_write_percent																																																										
1	2021-10-06 13:03:31.780	81.17	0.00	0.00																																																										
2	2021-10-06 13:03:16.770	77.06	0.00	0.00																																																										
3	2021-10-06 13:03:01.770	42.29	0.00	0.00																																																										
4	2021-10-06 13:02:46.747	0.00	0.00	0.00																																																										
5	2021-10-06 13:02:31.727	0.00	0.00	0.00																																																										
6	2021-10-06 13:02:16.680	0.00	0.00	0.00																																																										
7	2021-10-06 13:02:01.633	0.00	0.96	0.00																																																										
8	2021-10-06 13:01:46.600	0.00	0.00	0.00																																																										
9	2021-10-06 13:01:31.570	0.00	0.00	0.00																																																										

Narrative/Instructions	Screenshot	Notes
<p>Let's try now query 2 to visualize graphically the resource consumption.</p> <p>Execute Query 2 and on the results, activate the Chart view</p> <p>On the Right</p>  <p>Settings</p> <p>Chart Type Line</p> <p>Data Direction Vertical</p> <p>Data Type Number</p> <p><input checked="" type="checkbox"/> Use column names as labels</p>		<p>If more than 1 hour of statistics is needed, use the DMV [sys].[server_resource_stats] up to 14 days</p>

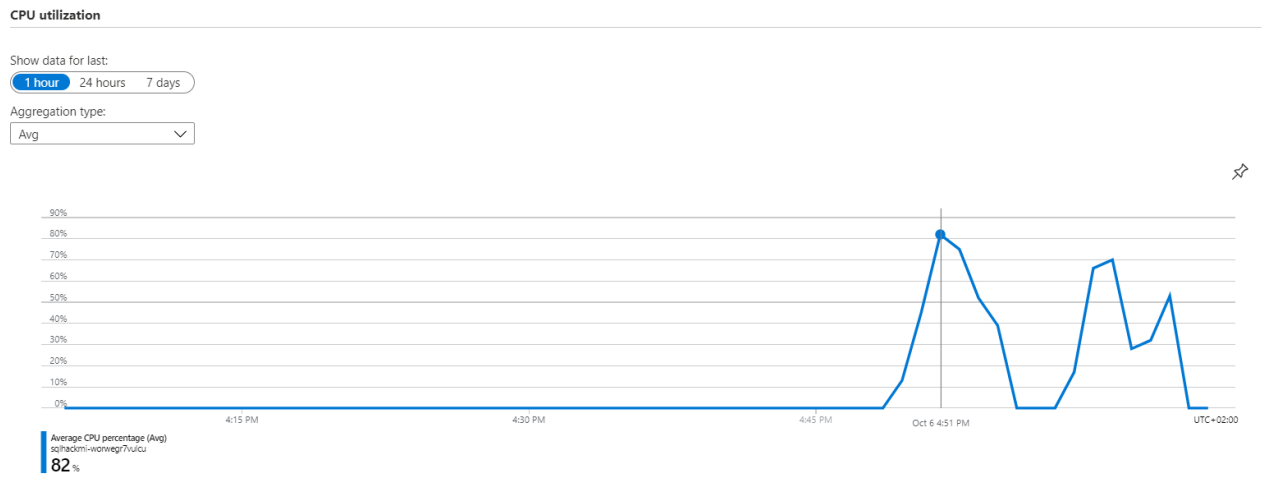
Instructions/Narrative	Screenshot	Notes																												
<p>Open SQL Server Management Studio using the same details:</p> <p>Server name: SQL MI Authentication: SQL Server Login: Will be provided during hack Password: Will be provided during hack</p>																														
<p>Open the SQL script: C:_SQLHACK_\LABS\02-Administering_Monitoring\Part_02_Monitoring_Lab_1.sql</p> <p>Run the PART 2 query 1. Identify the most consuming stored procedure based on the total_worker_time (CPU)</p>	<table><tr><th>ProcedureName</th><th>total_worker_time_ms</th><th>min_worker_time</th><th>max_worker_time_ms</th></tr><tr><td>uspGetManagerEmployees</td><td>538</td><td>0</td><td>26</td></tr><tr><td>uspGetWhereUsedProductID</td><td>561</td><td>1</td><td>4</td></tr><tr><td>GetTransactionTotalPerReferenceOrderID</td><td>3167850</td><td>647</td><td>1011</td></tr><tr><td>uspGetBillOfMaterials</td><td>113</td><td>0</td><td>2</td></tr><tr><td>sp_get_sqlagent_properties</td><td>94</td><td>30</td><td>32</td></tr><tr><td>uspGetEmployeeManagers</td><td>1351</td><td>2</td><td>8</td></tr></table>	ProcedureName	total_worker_time_ms	min_worker_time	max_worker_time_ms	uspGetManagerEmployees	538	0	26	uspGetWhereUsedProductID	561	1	4	GetTransactionTotalPerReferenceOrderID	3167850	647	1011	uspGetBillOfMaterials	113	0	2	sp_get_sqlagent_properties	94	30	32	uspGetEmployeeManagers	1351	2	8	
ProcedureName	total_worker_time_ms	min_worker_time	max_worker_time_ms																											
uspGetManagerEmployees	538	0	26																											
uspGetWhereUsedProductID	561	1	4																											
GetTransactionTotalPerReferenceOrderID	3167850	647	1011																											
uspGetBillOfMaterials	113	0	2																											
sp_get_sqlagent_properties	94	30	32																											
uspGetEmployeeManagers	1351	2	8																											
<p>Then to have further details on the most consuming stored procedure</p> <p>Run the PART 2 query 2, replace the filter clause with the stored procedure name. What is the issue? How to optimise?</p> <p>Tip: Have a look to the execution plan</p>	<p>CROSS APPLY and Table-valued Functions</p> <p>The PART 2 query joins the [sys].[dm_exec_requests] and [sys].[dm_exec_sql_text] DMVs to obtain long running batches and critically their offending SQL.</p> <p>Note that [dm_exec_sql_text] is actually a table-valued function (TVF) hence the use of the CROSS APPLY as TVFs can't be used in a normal join operation. The CROSS APPLY therefore produces the required inner-join between [dm_exec_requests] and [dm_exec_sql_text].</p> <p>See APPLY documentation here: Using APPLY Microsoft Docs and this excellent explanation on MSSQLTips: SQL Server CROSS APPLY and OUTER APPLY (mssqltips.com)</p>																													

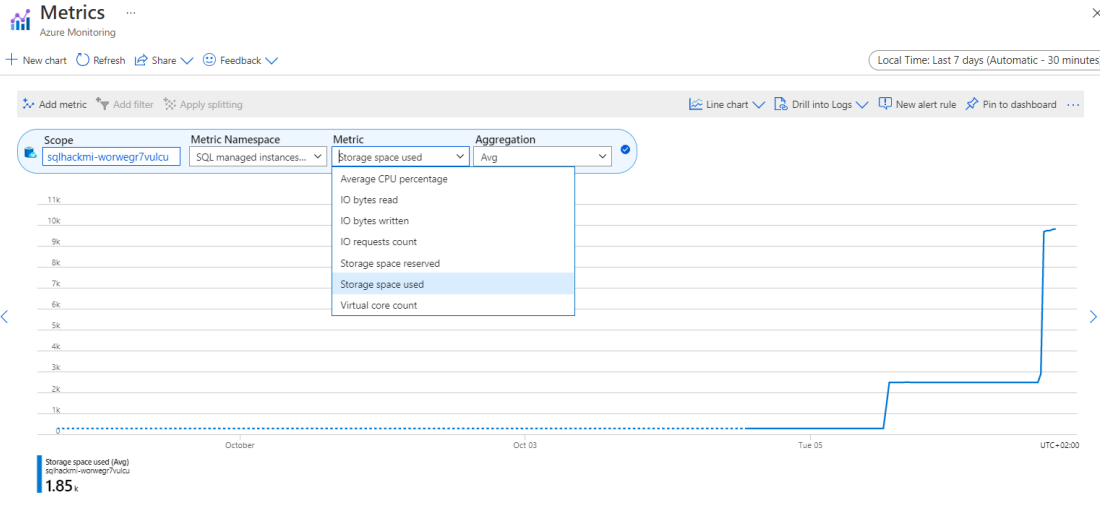
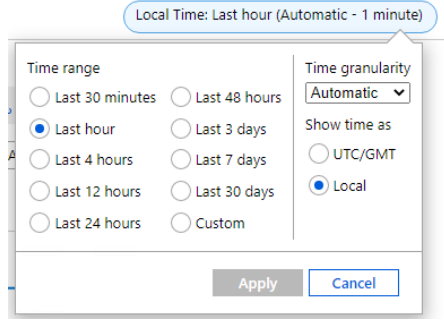
<p>In SSMS expand the [TenantCRM] database then expand Query Store.</p> <p>Remember that the Query Store collects telemetry data every 15 seconds and persists it for about 1hr.</p> <p>Double click on Overall Resource Consumption report to open it.</p>	 <ul style="list-style-type: none"> TenantCRM <ul style="list-style-type: none"> Database Diagrams Tables Views External Resources Synonyms Programmability Query Store <ul style="list-style-type: none"> Regressed Queries Overall Resource Consumption Top Resource Consuming Queries Queries With Forced Plans Queries With High Variation Query Wait Statistics Tracked Queries Service Broker Storage Security
--	--

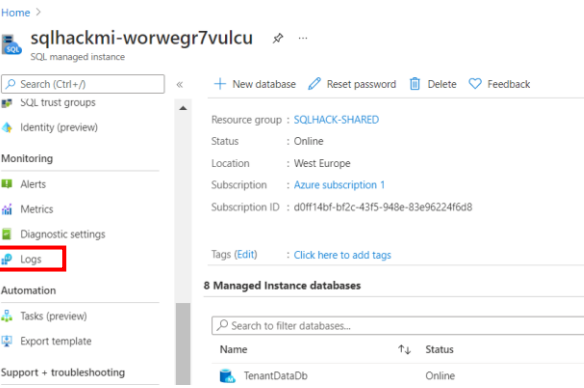
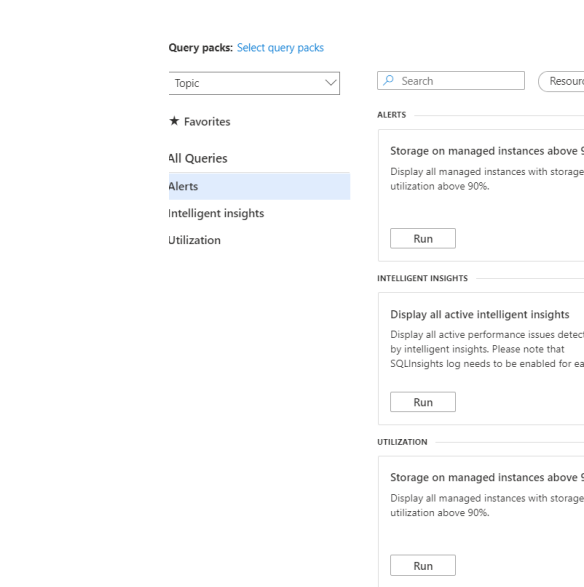
Narrative/Instructions	Screenshot	Notes
<p>For further details CPU Time, double click on the current day bar</p>		
<p>You can identify the top CPU consuming queries.</p> <p>Click on bars for more details</p>		

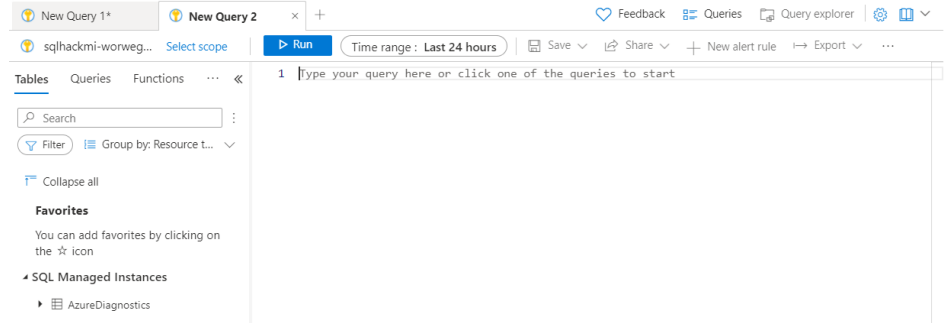
2. Using the Azure Portal to identify performance bottlenecks

As we are using Azure SQL Managed Instance we can also use the richness of the Azure Portal to help us to troubleshoot performance issues. Let us see how this is done.

Instructions/Narrative	Screenshot	Notes																												
<p>Open the Azure Portal (portal.azure.com) and open the shared SQL Managed Instance.</p> <p>Scroll down the Overview screen until you see the performance graph showing CPU utilization.</p> <p>Toggle the switches next to view CPU consumption over 1 hour, 24 hours and 7 days.</p> <p>Has the CPU usage pattern changed over time?</p>	 <p>The screenshot shows the 'CPU utilization' graph in the Azure Portal. The graph displays CPU percentage over time, with a peak around 82%. The x-axis shows time from 4:15 PM to UTC+02:00. The y-axis shows CPU percentage from 0% to 90%. The graph is titled 'CPU utilization' and includes controls for 'Show data for last:' (1 hour, 24 hours, 7 days) and 'Aggregation type:' (Avg). A legend at the bottom left indicates 'Average CPU percentage (Avg)' for 'sqlhackmi-vorweg7vulcu' with a value of 82%.</p> <table><caption>Approximate data points from the CPU utilization graph</caption><tr><th>Time</th><th>CPU Percentage (%)</th></tr><tr><td>4:15 PM</td><td>0</td></tr><tr><td>4:30 PM</td><td>0</td></tr><tr><td>4:45 PM</td><td>0</td></tr><tr><td>Oct 6 4:51 PM</td><td>82</td></tr><tr><td>5:00 PM</td><td>70</td></tr><tr><td>5:15 PM</td><td>40</td></tr><tr><td>5:30 PM</td><td>0</td></tr><tr><td>5:45 PM</td><td>0</td></tr><tr><td>6:00 PM</td><td>65</td></tr><tr><td>6:15 PM</td><td>70</td></tr><tr><td>6:30 PM</td><td>30</td></tr><tr><td>6:45 PM</td><td>50</td></tr><tr><td>7:00 PM</td><td>0</td></tr></table>	Time	CPU Percentage (%)	4:15 PM	0	4:30 PM	0	4:45 PM	0	Oct 6 4:51 PM	82	5:00 PM	70	5:15 PM	40	5:30 PM	0	5:45 PM	0	6:00 PM	65	6:15 PM	70	6:30 PM	30	6:45 PM	50	7:00 PM	0	
Time	CPU Percentage (%)																													
4:15 PM	0																													
4:30 PM	0																													
4:45 PM	0																													
Oct 6 4:51 PM	82																													
5:00 PM	70																													
5:15 PM	40																													
5:30 PM	0																													
5:45 PM	0																													
6:00 PM	65																													
6:15 PM	70																													
6:30 PM	30																													
6:45 PM	50																													
7:00 PM	0																													

Instructions/Narrative	Screenshot	Notes
<p>Click into the graph to go into Metrics chart screen. Here you have more flexibility over the data being reported and the format of the chart. Let's add Storage Space Used the chart.</p> <p>Click Add Metric</p> <p>In the Metric drop down select Storage Space Used.</p>		<p>Notice that from the chart screen you can alter the type of chart (Line, Area, Bar, Scatter to data Grid), set up alerts, and pin the chart to your Azure Portal dashboard.</p>
<p>Click Local Time on the top right of the chart to view the data for various time intervals.</p>		

Instructions/Narrative	Screenshot	Notes								
<p>On the Azure Portal, on the SQL Managed Instance Resource, click on Logs on the left Menu</p>	 <p>Home > sqlhackmi-worwegr7vulcu SQL managed instance</p> <p>Search (Ctrl+/) << + New database Reset password Delete Feedback</p> <p>SQL trust groups Identity (preview) Monitoring Alerts Metrics Diagnostic settings Logs Automation Tasks (preview) Export template Support + troubleshooting</p> <p>Resource group : SQLHACK-SHARED Status : Online Location : West Europe Subscription : Azure subscription 1 Subscription ID : d0ff14bf-bf2c-43f5-948e-83e96224f6d8</p> <p>Managed instance admin : DemoUser Host : sqlhackmi-worwegr7v Pricing tier : General Purpose Gen5 Instance pool : Not in an instance pool Virtual network / subnet : SQLHACK-SHARED-vn Virtual cluster : VirtualCluster364b040</p> <p>Tags (Edit) : Click here to add tags</p> <p>8 Managed Instance databases</p> <table><thead><tr><th>Name</th><th>Status</th><th>Creation time (UTC)</th><th>Earl</th></tr></thead><tbody><tr><td>TenantDataDb</td><td>Online</td><td>2021-10-05 10:36</td><td>2021</td></tr></tbody></table>	Name	Status	Creation time (UTC)	Earl	TenantDataDb	Online	2021-10-05 10:36	2021	
Name	Status	Creation time (UTC)	Earl							
TenantDataDb	Online	2021-10-05 10:36	2021							
<p>You will find several query templates</p> <p>You can test the CPU utilization threshold template</p> <p>And modify the parameters if needed.</p>	 <p>Query packs: Select query packs</p> <p>Topic Search Resource type : SQL Managed Instances Add filter</p> <p>★ Favorites All Queries Alerts Intelligent insights Utilization</p> <p>ALERTS</p> <p>Storage on managed instances above 9... Display all managed instances with storage utilization above 90%. Run Example query</p> <p>CPU utilization threshold above 95% on ... Display all managed instances with CPU threshold being over 95% of threshold. Run Example query</p> <p>INTELLIGENT INSIGHTS</p> <p>Display all active intelligent insights Intelligent insights report detecting the workload behavior as continuously hitting CPU limits. Please note that SQLInsights log needs to be enabled for each... Run Example query</p> <p>Workload continuously hitting CPU limits Intelligent insights report detecting the workload behavior as continuously hitting CPU limits. Please note that SQLInsights log needs to be enabled for each... Run Example query</p> <p>UTILIZATION</p> <p>Storage on managed instances above 9... Display all managed instances with storage utilization above 90%. Run Example query</p> <p>CPU utilization threshold above 95% on ... Display all managed instances with CPU threshold being over 95% of threshold. Run Example query</p>									

Instructions/Narrative	Screenshot	Notes
<p>Create a new Query</p>		
<p>You can try the following queries. AzureDiagnostics is the table receiving the different diagnostics</p> <p>Modify the Time Range, By default 24hours</p>	<pre>//general queries // query 1 without filter AzureDiagnostics // query 2 SQL MI Resource usage AzureDiagnostics where Resource == "SQLHACKMI-WORWEGR7VULCU" //replace with your SQL MI Instance Name where Category == "ResourceUsageStats" // query 3 on database: waits Stats AzureDiagnostics where Resource == "TENANTCRM" where Category == "QueryStoreWaitStatistics" // query 4 on database: errors AzureDiagnostics where Resource == "TENANTCRM" where Category == "Errors"</pre>	<p>Language Reference</p> <p>SQL to Kusto query translation - Azure Data Explorer Microsoft Docs</p>
<p>On the results for each of the queries, change the to Chart</p>	<pre>// more specific queries //Query get CPU usage AzureDiagnostics</pre>	


Results | Chart

The chart Formatting is on the right.

```
| where Resource == "SQLHACKMI-WORWEGR7VULCU" and TimeGenerated > ago(2h) //replace  
with your SQL MI Instance Name  
| where Category == "ResourceUsageStats"  
| project TimeGenerated, todecimal(avg_cpu_percent_s)  
// visualize the resulet as a chart
```

Results | Chart | ⌚ Display time (UTC+02:00) ▾

Completed 00:00.4 177 records



avg_cpu_percent_s

TimeGenerated [Local Time]

avg_cpu_percent_s

Chart formatting

Chart type

Line

▼ Data

X axis

TimeGenerated

Y axis

avg_cpu_percent_s

Split-by

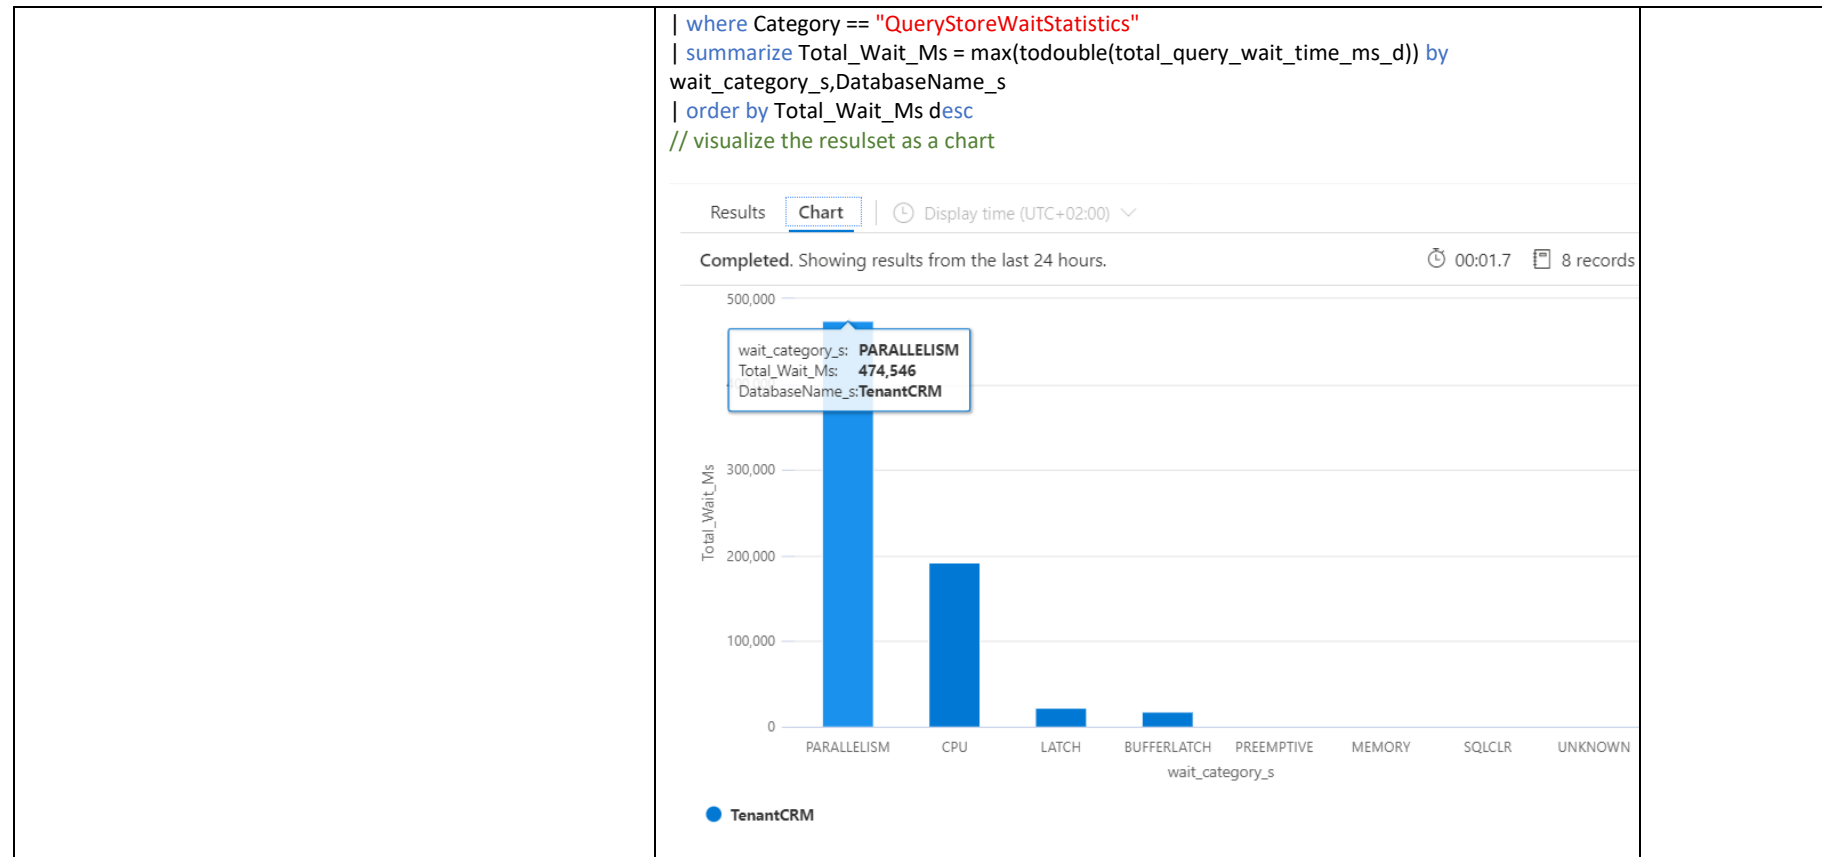
Select a column

Aggregation

Sum

> Legend

```
// Query get database Waits Stats  
AzureDiagnostics
```



We've now seen how the various performance monitoring and diagnosis tools – the Azure Portal, DMVs and the Query Store reports all revealed that the CPU was under pressure. The DMVs and Query Store reports also allowed us to drill-down to root cause and identify the worst offending query. In the real world you would now progress to tune the offending query to reduce the overall load on the Managed Instances CPUs and thereby improve the databases and environment performance.